

Diffusion Models & Inpainting

DDP - DDIM - DPS - CFG

Angela Saade - Aurélien Daudin - Baptiste Arnold
Khaled Mili - Maxime Ruff - Pierre Schweitzer

EPITA

January 26, 2026

Outline

1 DDPM

- Introduction
- Simplified Training Objective
- Limitations

2 DDIM

- Formulation
- Conclusion

3 DPS - Diffusion Posterior Sampling

- Introduction
- Posterior Distribution and Likelihood
- Approximation

DDPM - Denoising Diffusion Probabilistic Models

The DDPM framework consists of two opposing Markov chains.

1. The Forward Process (q):

- A fixed chain that gradually adds Gaussian noise according to a variance schedule β_t .
- Transition kernel:

$$q(x_t | x_{t-1}) := \mathcal{N}(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

2. The Reverse Process (p_θ):

- A learned Markov chain with Gaussian transitions.
- Parameterized by a neural network to reverse the noise:

$$p_\theta(x_{t-1} | x_t) := \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

DDPM - Efficient Forward Sampling

A key property allows us to sample x_t at any timestep t directly from x_0 without iterating.

Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$.

Marginal at timestep t

$$q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$$

Practical Implication: We can generate noisy training samples on the fly:

$$x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, \quad \text{where } \epsilon \sim \mathcal{N}(0, I)$$

This makes training efficient by optimizing random terms of the variational bound.

DDPM - Simplified Training Objective (L_{simple})

Instead of predicting the mean $\tilde{\mu}_t$, the network ϵ_θ is trained to predict the **noise** ϵ added to the image.

The DDPM paper proposes a simplified weighted variational bound:

Loss Function

$$L_{simple}(\theta) := \mathbb{E}_{t, x_0, \epsilon} \left[\|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2 \right]$$

Key Insights:

- It resembles denoising score matching over multiple noise scales.
- It down-weights loss terms at small t to focus on difficult denoising tasks.

Sampling with Langevin Dynamics

To generate data, we start from pure noise $x_T \sim \mathcal{N}(0, I)$ and iterate backwards.

Sampling Step

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) + \sigma_t z$$

- $\epsilon_\theta(x_t, t)$ estimates the noise to be removed.
- $z \sim \mathcal{N}(0, I)$ adds stochasticity (Langevin dynamics).
- This stochastic term prevents the process from becoming deterministic and collapsing to the mean.

DDPM Limitations

The Problem with DDPM:

- The generative process approximates the reverse of a diffusion process.
- **Consequence:** Sampling is extremely slow. Generating a single image requires running the neural network 1000 times .

The DDIM Insight:

- The training objective L_{simple} only depends on the marginals $q(x_t|x_0)$, not the specific joint trajectory.
- We can define a **non-Markovian forward process** that leads to the exact same marginals but allows for a different, deterministic reverse process .

Denoising Diffusion Implicit Models (DDIM)

DDIM generalizes the sampling equation.

New Update Equation (Eq. 12 in paper): To go from x_t to x_{t-1} , we predict x_0 (denoted as $f_\theta^{(t)}(x_t)$) and then interpolate:

Generic DDIM Sampling Step

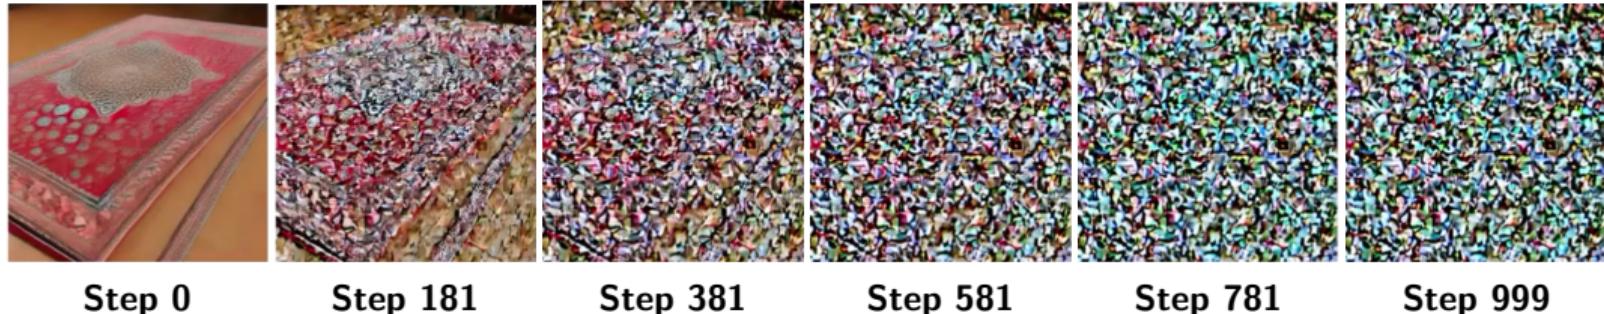
$$x_{t-1} = \underbrace{\sqrt{\bar{\alpha}_{t-1}} \left(\frac{x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_\theta(x_t)}{\sqrt{\bar{\alpha}_t}} \right)}_{\text{Denoised } x_0 \text{ prediction}} + \underbrace{\sqrt{1 - \bar{\alpha}_{t-1} - \sigma_t^2} \cdot \epsilon_\theta(x_t)}_{\text{Direction pointing to } x_t} + \underbrace{\sigma_t \epsilon}_{\text{Random noise}}$$

The DDIM Case ($\sigma_t = 0$):

- By setting $\sigma_t = 0$, the random noise term vanishes.
- The process becomes **deterministic** (Implicit Model).
- We can use the **same pre-trained model** ϵ_θ as DDPM!.

DDIM - Visuals

$$\mathbf{x}_{t-1} = \underbrace{\sqrt{\alpha_{t-1}} \left(\frac{\mathbf{x}_t - \sqrt{1-\alpha_t} \epsilon_\theta^{(t)}(\mathbf{x}_t)}{\sqrt{\alpha_t}} \right)}_{\text{"predicted } \mathbf{x}_0\text{"}} + \boxed{\underbrace{\sqrt{1-\alpha_{t-1}-\sigma_t^2} \cdot \epsilon_\theta^{(t)}(\mathbf{x}_t) + \underbrace{\sigma_t \epsilon_t}_{\text{random noise}}}_{\text{"direction pointing to } \mathbf{x}_t\text{"}}}$$



Step 0

Step 181

Step 381

Step 581

Step 781

Step 999

DDPM vs. DDIM: Conclusion

Despite using the **same neural network** trained with the **same loss function**, their inference behaviors differ significantly.

Feature	DDPM	DDIM
Reverse Process	Stochastic (Markovian).	Deterministic (Non-Markovian).
Speed	Slow (requires ~ 1000 steps).	Fast (high quality in 50-100 steps).
Latent Space (x_T)	No strong link to x_0 (stochastic mapping).	Fixed encoding of x_0 (allows interpolation and reconstruction).

Classifier Guidance (CG): Mathematical Framework

Goal: Modify conditional score by incorporating signal from auxiliary classifier

Starting point (standard conditional score):

$$\epsilon_\theta(x_t, c) \approx -\sigma_t \nabla_{x_t} \log p(x_t | c) \quad (1)$$

Add classifier gradient with weight w :

$$\tilde{\epsilon}(x_t, c) = \epsilon_\theta(x_t, c) - w \sigma_t \nabla_{x_t} \log p_\phi(c | x_t) \quad (2)$$

Implicit distribution: Sampling with this modified score approximates

$$\tilde{p}_\theta(x_t | c) \propto p_\theta(x_t | c) \cdot [p_\phi(c | x_t)]^w \quad (3)$$

Intuition:

- Base term: $\epsilon_\theta(x_t, c)$ predicts natural denoising direction
- Correction term: $-w \sigma_t \nabla \log p_\phi$ pushes toward regions where classifier is confident
- Weight w controls strength of this correction

Classifier Guidance: Complete Mathematical Derivation

Step 1: Bayes' rule for conditional probability

$$\nabla_{x_t} \log p(x_t | c) = \nabla_{x_t} \log p(x_t) + \nabla_{x_t} \log p(c | x_t) \quad (4)$$

Step 2: Scale both sides by $(1 + w)$

$$(1 + w) \nabla_{x_t} \log p(x_t | c) = (1 + w) \nabla_{x_t} \log p(x_t) \quad (5)$$

$$+ (1 + w) \nabla_{x_t} \log p(c | x_t) \quad (6)$$

Step 3: Expand and rearrange

$$= \nabla_{x_t} \log p(x_t) + w \nabla_{x_t} \log p(x_t) + \nabla_{x_t} \log p(x_t | c) \quad (7)$$

$$+ w \nabla_{x_t} \log p(c | x_t) \quad (8)$$

Step 4: Convert to noise predictors using $\epsilon \approx -\sigma \nabla \log p$

Result:

$$\tilde{\epsilon}(x_t, c) = \epsilon_\theta(x_t, c) - w \sigma_t \nabla_{x_t} \log p_\phi(c | x_t) \quad (9)$$

Classifier Guidance: Inference Algorithm

Algorithm 1 Classifier Guidance Sampling

Require: condition c , guidance strength w , timesteps $\{t_1, \dots, t_T\}$

```
1: Initialize:  $x_T \sim \mathcal{N}(0, I)$ 
2: for  $t = T$  to 1 do
3:    $\epsilon_{\text{model}} \leftarrow \epsilon_{\theta}(x_t, t, c)$     // Diffusion forward
4:   Compute  $\nabla_{x_t} \log p_{\phi}(c | x_t)$     // Classifier forward + backward
5:    $\tilde{\epsilon} \leftarrow \epsilon_{\text{model}} - w\sigma_t \nabla_{x_t} \log p_{\phi}(c | x_t)$     // Guidance
6:    $x_{t-1} \sim p_{\theta}(x_{t-1} | x_t, c, \tilde{\epsilon})$     // Update step
7: end for
8: return  $x_0 = x_1$ 
```

Computational requirements per timestep:

- ✓ 1 forward pass: diffusion model $\epsilon_{\theta}(x_t, t, c)$
- ✓ 1 forward pass: classifier $p_{\phi}(c | x_t)$
- ✓ 1 backward pass: classifier gradient $\nabla_{x_t} \log p_{\phi}(c | x_t)$
- ✓ 1 linear combination (negligible)

Classifier-Free Guidance (CFG): Key Innovation

Central idea: Use implicit classifier derived from the generative model itself

Implicit classifier (by Bayes' rule):

By Bayes' theorem, we can express the class probability as:

$$p_i(c | x_t) = \frac{p(x_t | c) p(c)}{p(x_t)} \propto \frac{p(x_t | c)}{p(x_t)} \quad (10)$$

Classifier-Free Guidance: Complete Derivation

Step 2: Gradient of log (score of implicit classifier)

$$\nabla_{x_t} \log p_i(c | x_t) = \nabla_{x_t} \log p(x_t | c) - \nabla_{x_t} \log p(x_t) \quad (11)$$

Step 3: Apply score-noise relationship

$$= -\frac{1}{\sigma_t} [\epsilon_\theta(x_t, c) - \epsilon_\theta(x_t)] \quad (12)$$

Step 4: Substitute into CG formula

$$\tilde{\epsilon}(x_t, c) = \epsilon_\theta(x_t, c) - w\sigma_t \cdot \left(-\frac{1}{\sigma_t}\right) [\epsilon_\theta(x_t, c) - \epsilon_\theta(x_t)] \quad (13)$$

$$\boxed{\tilde{\epsilon}(x_t, c) = (1 + w)\epsilon_\theta(x_t, c) - w\epsilon_\theta(x_t)} \quad (14)$$

Classifier-Free Guidance: Training Strategy

Key idea: Train **ONE U-Net** for both conditional and unconditional generation

Training algorithm:

Algorithm 2 Joint Training for CFG

Require: training data (x, c) , unconditional dropout probability p_{uncond}

```
1: repeat
2:   for each batch  $(x, c)$  do
3:     With probability  $p_{\text{uncond}}$ : set  $c \leftarrow \emptyset$  (null token)
4:     Sample timestep:  $t \sim \text{Uniform}(1, T)$ 
5:     Sample noise:  $\epsilon \sim \mathcal{N}(0, I)$ 
6:     Corrupt:  $x_t = \sqrt{\alpha_t}x + \sqrt{1 - \alpha_t}\epsilon$ 
7:     Compute loss:  $\mathcal{L} = \|\epsilon_\theta(x_t, t, c) - \epsilon\|_2^2$ 
8:     Update:  $\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}$ 
9:   end for
10:  until converged
```

Classifier-Free Guidance: Inference Algorithm

Algorithm 3 Classifier-Free Guidance Sampling

Require: condition c , guidance strength w , timesteps $\{t_1, \dots, t_T\}$

- 1: **Initialize:** $x_T \sim \mathcal{N}(0, I)$
 - 2: **for** $t = T$ to 1 **do**
 - 3: $\epsilon_{\text{cond}} \leftarrow \epsilon_\theta(x_t, t, c)$ // Conditional prediction
 - 4: $\epsilon_{\text{uncond}} \leftarrow \epsilon_\theta(x_t, t, \emptyset)$ // Unconditional prediction
 - 5: $\tilde{\epsilon} \leftarrow (1 + w)\epsilon_{\text{cond}} - w\epsilon_{\text{uncond}}$ // Linear combination
 - 6: $x_{t-1} \sim p_\theta(x_{t-1} | x_t, c, \tilde{\epsilon})$ // Update step
 - 7: **end for**
 - 8: **return** $x_0 = x_1$
-

Computational requirements per timestep:

- ✓ 2 forward passes: same U-Net (conditional + unconditional)
- ✗ 0 classifier models
- ✗ 0 gradient computations
- ✓ 1 linear combination (trivial)

Comparison: CG vs CFG + Practical Guidance Values

Aspect	Classifier Guidance (CG)	Classifier-Free (CFG)
Inference cost	2 forwards + 1 backward	2 forwards
Models needed	Diffusion + Classifier	Diffusion only
Training complexity	High (classifier on noisy)	Low (random dropout)
Peak metrics (FID/IS)	Higher	Competitive
Industry adoption	Limited	Standard (Stable Diffusion, SDXL)
Artifacts at high w	Adversarial edges	Color saturation

Inverse Problems with Diffusion

The Goal: Recover clean image x from degraded measurement $y = f(x) + n$

Challenge: Standard DDPM samples from prior $p(x)$ (random images)

Requirement: Sample from posterior $p(x|y)$ (conditional on measurement)

DDPM Limitation:

- ✗ Ignores the observation y
- ✗ Generates images unrelated to input
- ✓ Solution: Add **data fidelity constraint**

Combining Scores with Bayes' Rule

We want to sample from posterior $p(x_t|y)$ instead of prior $p(x_t)$.

Bayes' theorem applied to score:

$$\nabla_{x_t} \log p(x_t|y) = \underbrace{\nabla_{x_t} \log p(x_t)}_{\text{Prior}} + \underbrace{\nabla_{x_t} \log p(y|x_t)}_{\text{Likelihood}} \quad (15)$$

Two terms:

- **Prior term:** $\nabla_{x_t} \log p(x_t)$ — Given by U-Net $\epsilon_\theta(x_t)$
 - Says: "This direction is realistic"
- **Likelihood term:** $\nabla_{x_t} \log p(y|x_t)$ — **Intractable!**
 - Says: "This direction respects measurement y "
 - Problem: x_t is noisy; need to integrate over all possible clean x_0

Diffusion Posterior Sampling (DPS): The Approximation

Key Idea: Use the model's current denoising estimate to approximate likelihood

At each step t , the U-Net denoising x_t implicitly predicts clean image: $\hat{x}_0(x_t)$ (Tweedie estimate)

Approximate the likelihood:

$$p(y|x_t) \approx p(y | \hat{x}_0(x_t)) \quad (16)$$

Computing the likelihood gradient:

- ① Calculate measurement error: $\mathcal{D} = \|y - f(\hat{x}_0)\|^2$
- ② Compute gradient via backpropagation:

$$\nabla_{x_t} \log p(y|x_t) \approx -\zeta \cdot \nabla_{x_t} \mathcal{D} \quad (17)$$

Requires: Backpropagation through the U-Net $\rightarrow x_t$ becomes differentiable variable

Applications & Conclusion

Key Takeaways:

- ① **Zero-shot:** No retraining, define f at inference
- ② **General:** Linear, non-linear, any measurement operator
- ③ **Robust:** Soft constraints handle noisy measurements naturally
- ④ **Cost:** Expensive, but generalizes where specialists fail

References I

-  Dhariwal, P., & Nichol, A. (2021).
“Diffusion Models Beat GANs on Image Synthesis.”
arXiv preprint arXiv:2105.05233.
-  Ho, J., & Salimans, T. (2021).
“Classifier-Free Diffusion Guidance.”
NeurIPS 2021 Workshop on Deep Generative Models and Downstream Applications.
-  Ho, J., Jain, A., & Abbeel, P. (2020).
“Denoising Diffusion Probabilistic Models.”
Advances in Neural Information Processing Systems (NeurIPS).
-  Song, Y., Garg, S., Gao, J., Lindenbaum, O., Tripathi, A., & Ermon, S. (2021).
“Solving Inverse Problems Deterministically without Diffusion.”
arXiv preprint arXiv:2108.08713.